

# **CACSC19: AI Hardware and Tools Workshop**



**Name: Chaitanya Yadav**  
**Roll No: 2022UCA1849**  
**Class: CSAI-1**

## Problem Statement

Traditional irrigation methods often waste water by following fixed schedules without accounting for real-time environmental conditions. This can lead to over-irrigation (wasting water) or under-irrigation (harming crops), especially in regions facing climate uncertainty or water scarcity.

### Goal:

Develop a smart irrigation scheduling system that:

- Predicts soil moisture conditions using environmental sensor data.
- Recommends when and how much to irrigate.
- Supports deployment on low-power microcontrollers using TinyML.

### Data Source

The dataset titled `Irrigation Scheduling.csv` contains various sensor readings including:

- Temperature
- Humidity
- Soil moisture
- Rainfall
- Altitude
- Other geographical features

### Target Variable: `class`

Represents irrigation need with 4 labels:

- "Very Dry"
- "Dry"
- "Wet"
- "Very Wet"

### Tech Stack & Tools Used

- Programming Languages: Python, R
- ML/DL Frameworks: PyTorch, scikit-learn
- Model Optimization: Knowledge Distillation (Teacher-Student architecture)
- Model Format: ONNX
- Visualization Tools: Power BI, Streamlit
- Data Streaming & Processing: PySpark
- DevOps: Docker
- Microcontroller Platform: ESP32 (simulated on Wokwi)

## **Project Workflow**

### Data Preprocessing & Analysis

- Initial exploration and cleaning are done using R.
- Scikit-learn is used for encoding and standardizing inputs.

### Model Training

- A Teacher Model (deep neural network) is trained for high accuracy.
- A Student Model (lightweight DNN) is trained using knowledge distillation.

### Model Export

- The student model is exported to ONNX format for hardware deployment.

### Frontend Interface

- A Streamlit app visualizes predictions and model outputs interactively.
- Power BI is used to present insights, trends, and analytics from the data.

### Real-time Streaming & Analytics

- PySpark is used for large-scale analysis and ingestion.

### DevOps & Deployment

- Everything is containerized and versioned using Docker for reproducibility and deployment.

## Hardware Interface

- The exported ONNX model runs on an ESP32 via Wokwi simulation.
- Real-time sensor inputs are used to make predictions and activate irrigation components.

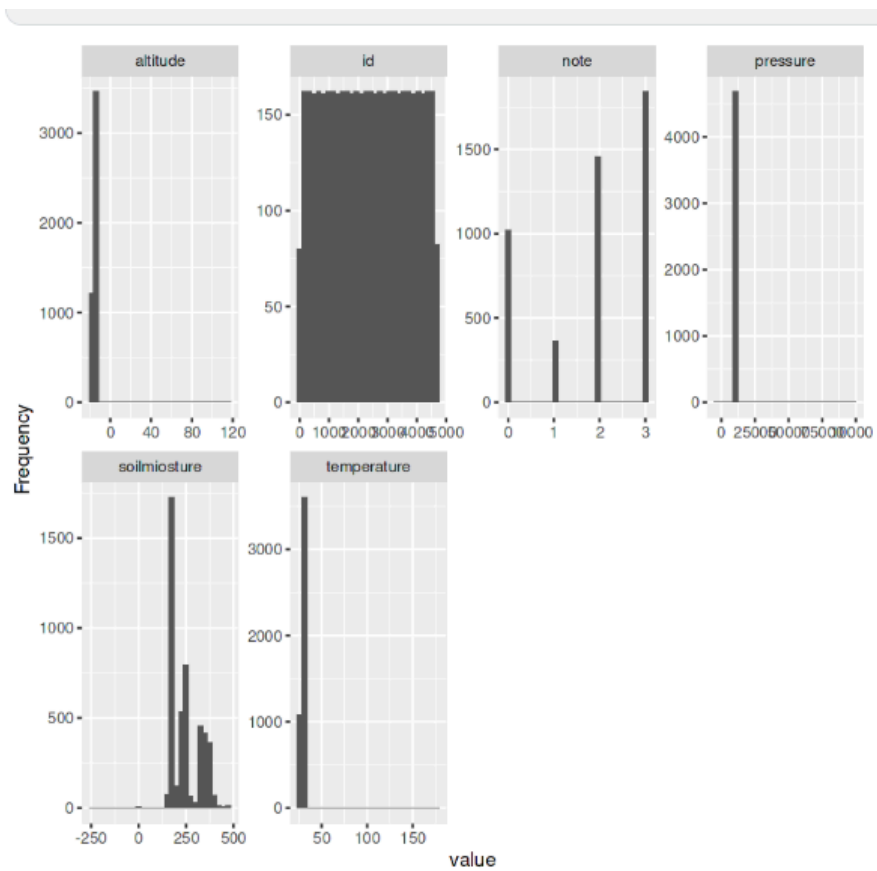
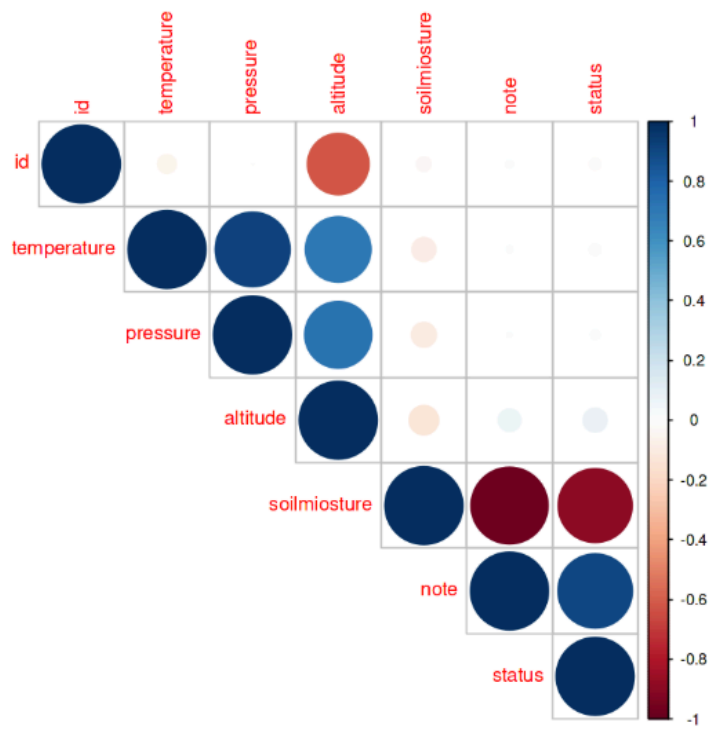
## 1. Data Preprocessing & Analysis

```
# Plot histograms for numeric variables
plot_histogram(data)

# Bar plots for categorical variables (if any)
plot_bar(data)
```

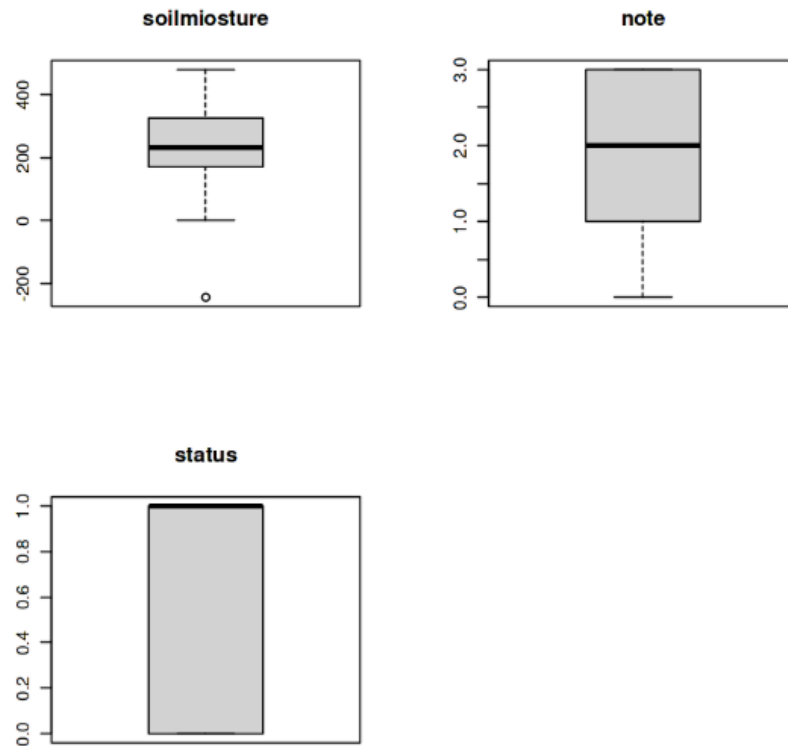
```
# Correlation matrix for numeric variables
numeric_data <- select_if(data, is.numeric)
cor_matrix <- cor(numeric_data, use = "complete.obs")
corrplot(cor_matrix, method = "circle", type = "upper")

# Pair plots
ggpairs(numeric_data)
```





```
# Scatter plots for selected relationships
ggplot(data, aes(x = data[[1]], y = data[[2]])) +
  geom_point() +
  labs(title = paste(names(data)[1], "vs", names(data)[2]))
```



## 2. Model Training

Windsurf: Refactor | Explain

```
class TeacherNet(nn.Module):
```

Windsurf: Refactor | Explain | Generate Docstring | X

```
def __init__(self, input_size, num_classes):
    super(TeacherNet, self).__init__()
    self.fc1 = nn.Linear(input_size, 128)
    self.fc2 = nn.Linear(128, 64)
    self.fc3 = nn.Linear(64, 32)
    self.out = nn.Linear(32, num_classes)
```

Windsurf: Refactor | Explain | Generate Docstring | X

```
def forward(self, x):
    x = F.relu(self.fc1(x))
    x = F.relu(self.fc2(x))
    x = F.relu(self.fc3(x))
    return self.out(x)
```

```

# Define student model
Windsurf: Refactor | Explain
class StudentNet(nn.Module):
    Windsurf: Refactor | Explain | Generate Docstring | ✕
    def __init__(self, input_size, num_classes):
        super(StudentNet, self).__init__()
        self.fc1 = nn.Linear(input_size, 32)
        self.out = nn.Linear(32, num_classes)

    Windsurf: Refactor | Explain | Generate Docstring | ✕
    def forward(self, x):
        x = F.relu(self.fc1(x))
        return self.out(x)

```

```

# Distillation training
student = StudentNet(X_train.shape[1], len(np.unique(y)))
# teacher = torch.load("teacher_model.pt")

import torch
from torch.serialization import add_safe_globals

# Register the class if it's a custom class
add_safe_globals([TeacherNet])

teacher = torch.load("teacher_model.pt", weights_only=False)
opt = torch.optim.Adam(student.parameters(), lr=0.001)

T, alpha = 3.0, 0.7

Windsurf: Refactor | Explain | Generate Docstring | ✕
def distill_loss(s_logit, t_logit, labels, T, alpha):
    h = F.cross_entropy(s_logit, labels)
    t_soft = F.log_softmax(t_logit / T, dim=1)
    s_soft = F.log_softmax(s_logit / T, dim=1)
    s_loss = F.kl_div(s_soft, t_soft, log_target=True, reduction='batchmean') * (T**2)
    return alpha * s_loss + (1 - alpha) * h

for epoch in range(20):
    student.train()
    teacher.eval()
    for xb, yb in train_loader:
        with torch.no_grad():
            t_logit = teacher(xb)
        s_logit = student(xb)
        loss = distill_loss(s_logit, t_logit, yb, T, alpha)
        opt.zero_grad(); loss.backward(); opt.step()

# Save final model
torch.save(student.state_dict(), "student_model.pt")

```

Do you want to save this file?  
REditor





AIHT-FINALPROJECT

- > \_\_pycache\_\_
- > venv
- app.py
- converter.py
- Dockerfile
- eda-aihtproject.ipynb
- frontend.py
- Irrigation Scheduling.csv
- model\_train\_and\_export.py
- README.md
- requirements.txt
- student\_model.onnx
- student\_model.pt
- teacher\_model.pt
- teacher.py
- temp.py
- test\_input.txt

```
from fastapi import FastAPI, File
import numpy as np
import onnxruntime as ort

app = FastAPI()
session = ort.InferenceSession("student_model.onnx")

Windsurf: Refactor | Explain | Generate Docstring | ✕
@app.post("/predict")
def predict_from_file(file: bytes = File(...)):
    content = file.decode("utf-8").strip()
    features = np.array([list(map(float, content.split(',')))], dtype=np.float32)
    outputs = session.run(["output"], {"input": features})
    predicted_class = int(np.argmax(outputs[0]))
    return {"predicted_class": predicted_class}
```

## 4. Frontend

```
import streamlit as st
import numpy as np
import onnxruntime as ort

# Load ONNX model
session = ort.InferenceSession("student_model.onnx")

st.title("🌱 Irrigation Scheduling Predictor")
st.write("Enter sensor readings below to predict irrigation class.")

feature_names = ["temperature", "humidity", "soil_moisture", "altitude", "rainfall", "wind_speed"]

inputs = []
for name in feature_names:
    value = st.number_input(f"{name}", value=0.0)
    inputs.append(value)

if st.button("Predict"):
    features = np.array([inputs], dtype=np.float32)
    output = session.run(["output"], {"input": features})
    predicted_class = int(np.argmax(output[0]))
    st.success(f"Predicted Irrigation Class: {predicted_class}")
```



## Irrigation Scheduling Predictor

Enter sensor readings below to predict irrigation class.

temperature

25.00

- +

humidity

2.00

- +

soil\_moisture

3.00

- +

altitude

15.00

- +

rainfall

9.00

- +

wind\_speed

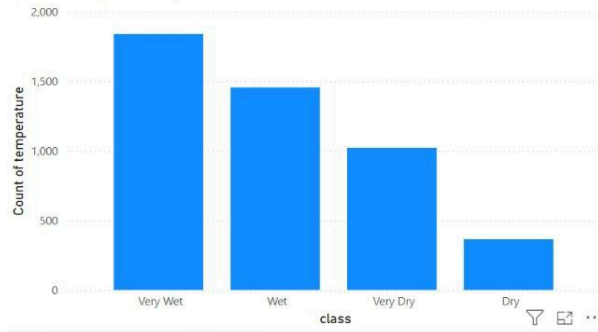
2.00

- +

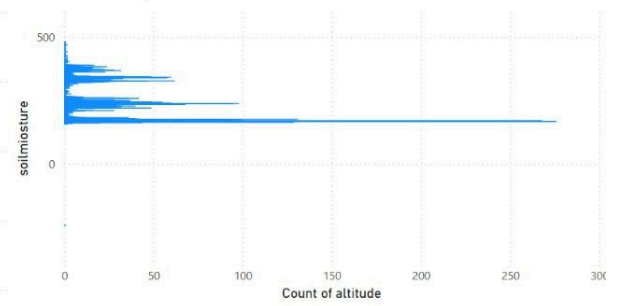
Predict

Predicted Irrigation Class: 3

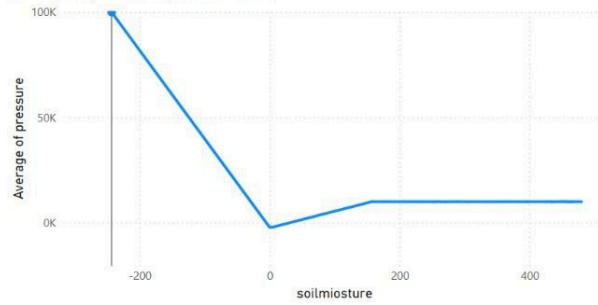
Count of temperature by class



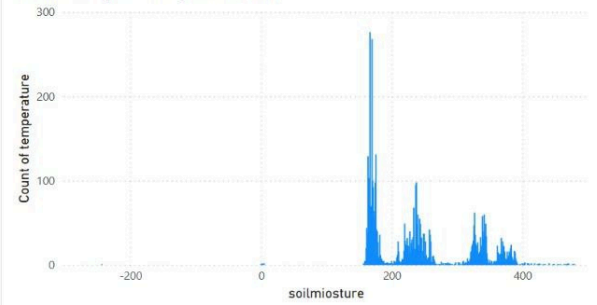
Count of altitude by soilmoisture



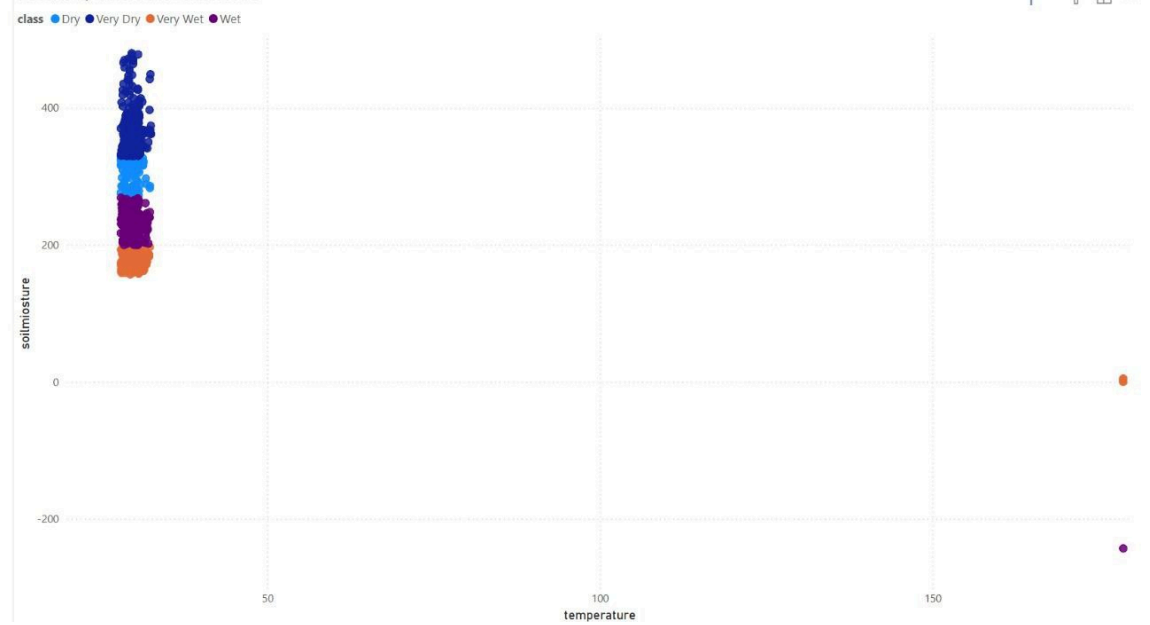
Average of pressure by soilmoisture

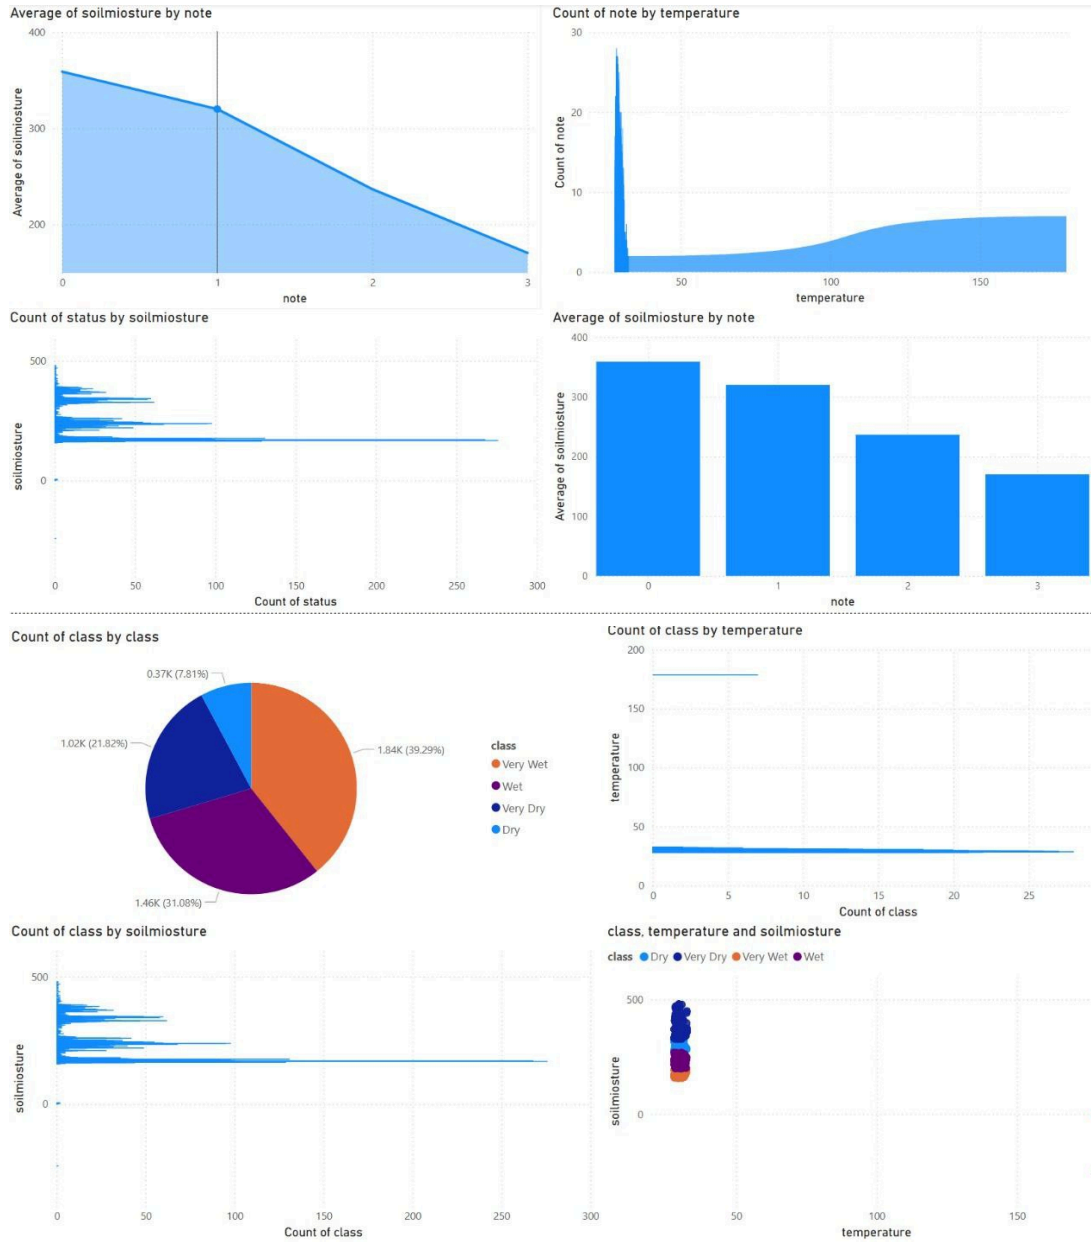


Count of temperature by soilmoisture



class, temperature and soilmoisture





## 5. Real-time Streaming

```

spark_teacher.py / ...
from pyspark.sql import SparkSession
from pyspark.sql.functions import col
from sklearn.preprocessing import LabelEncoder, StandardScaler
import pandas as pd
import numpy as np

# Initialize Spark
spark = SparkSession.builder.appName("IrrigationPipeline").getOrCreate()

# Load dataset
df_spark = spark.read.csv("Irrigation Scheduling.csv", header=True, inferSchema=True)

# Drop unnecessary columns and handle nulls
df_spark = df_spark.drop("id", "date", "time")
df_spark = df_spark.na.fill({"altitude": df_spark.select("altitude").agg({"altitude": "mean"}).first()[0]})

# Convert to Pandas for label encoding and scaling
df = df_spark.toPandas()

# Encode labels
le = LabelEncoder()
df["class_encoded"] = le.fit_transform(df["class"])

# Feature scaling
X = df.drop(columns=["class", "class_encoded"])
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
y = df["class_encoded"]

# Create final DataFrame
df_final = pd.DataFrame(X_scaled, columns=X.columns)
df_final["label"] = y

# Save scaled test set to simulate Kafka stream
train_df = df_final.sample(frac=0.8, random_state=42)
test_df = df_final.drop(train_df.index)

```

## 6. Deployment and DevOps

```

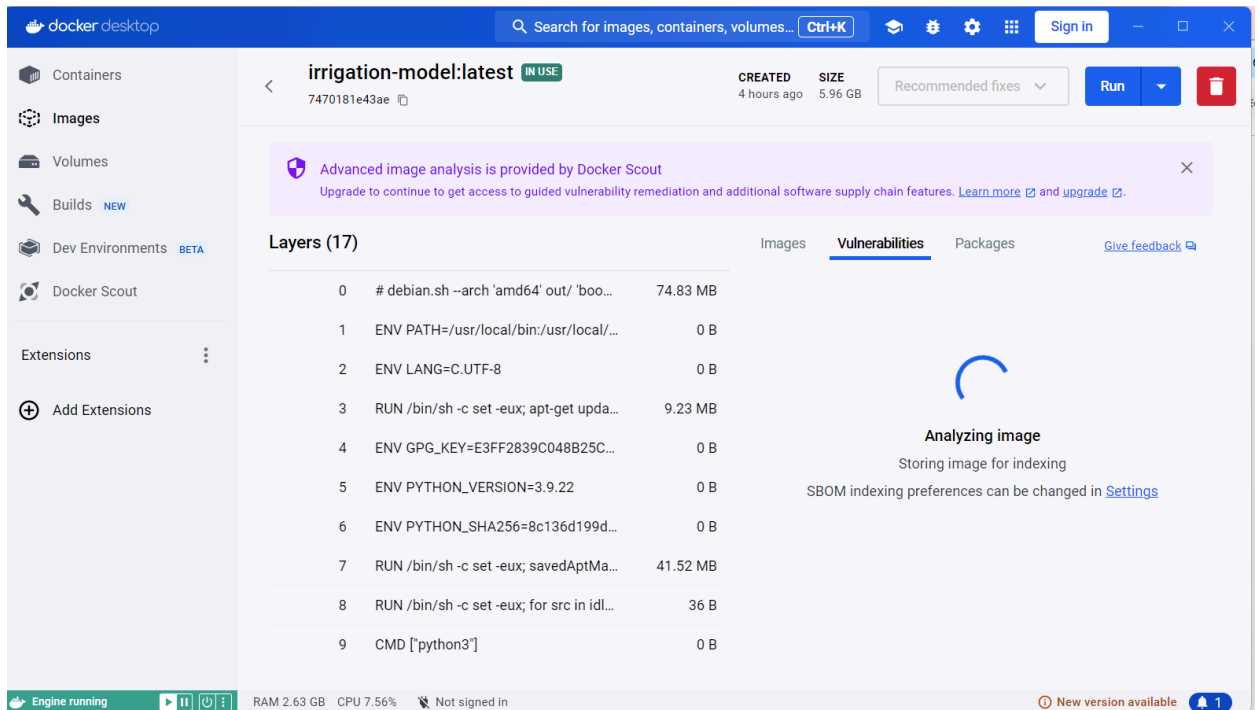
FROM python:3.9-slim
WORKDIR /app

COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

COPY app.py .
COPY student_model.onnx .

EXPOSE 8000
CMD ["uvicorn", "app:app", "--host", "0.0.0.0", "--port", "8000"]

```



## 7. Hardware Interface

```

!pip install Flask onnxruntime numpy

import flask
import onnxruntime as ort
import numpy as np
import os
import csv
from datetime import datetime
import threading

app = flask.Flask(__name__)

# --- Configuration ---
MODEL_PATH = 'student_model.onnx' # Make sure this file is deployed with your app
INPUT_NAME = None
OUTPUT_NAME = None
EXPECTED_INPUT_SHAPE = (1, 6)
LOG_FILE = 'data_log.csv' # <-- CSV log file
# --- End Configuration ---

session = None

def load_model():
    global session, INPUT_NAME, OUTPUT_NAME
    if not os.path.exists(MODEL_PATH):
        print(f"ERROR: Model file not found at {MODEL_PATH}")
        return False
    try:
        print(f"Loading ONNX model from {MODEL_PATH}...")
        session = ort.InferenceSession(MODEL_PATH, providers=['CPUExecutionProvider'])
        INPUT_NAME = session.get_inputs()[0].name
        OUTPUT_NAME = session.get_outputs()[0].name
        print(f"Model loaded successfully.")
        print(f"Input Name: {INPUT_NAME}, Expected Shape: {EXPECTED_INPUT_SHAPE}")
        print(f"Output Name: {OUTPUT_NAME}")
        return True
    
```



```

@app.route('/infer', methods=['POST'])
def infer():
    if session is None:
        print("Model not loaded, attempting reload...")
        if not load_model():
            return flask.jsonify({"error": "Model not loaded on server"}), 500

    try:
        data = flask.request.get_json()
        if not data or 'inputs' not in data:
            return flask.jsonify({"error": "Missing 'inputs' key in JSON payload"}), 400

        input_values = data['inputs']

        if not isinstance(input_values, list) or len(input_values) != EXPECTED_INPUT_SHAPE[1]:
            return flask.jsonify({"error": f"Expected a list of {EXPECTED_INPUT_SHAPE[1]} float values in 'inputs'"}), 400

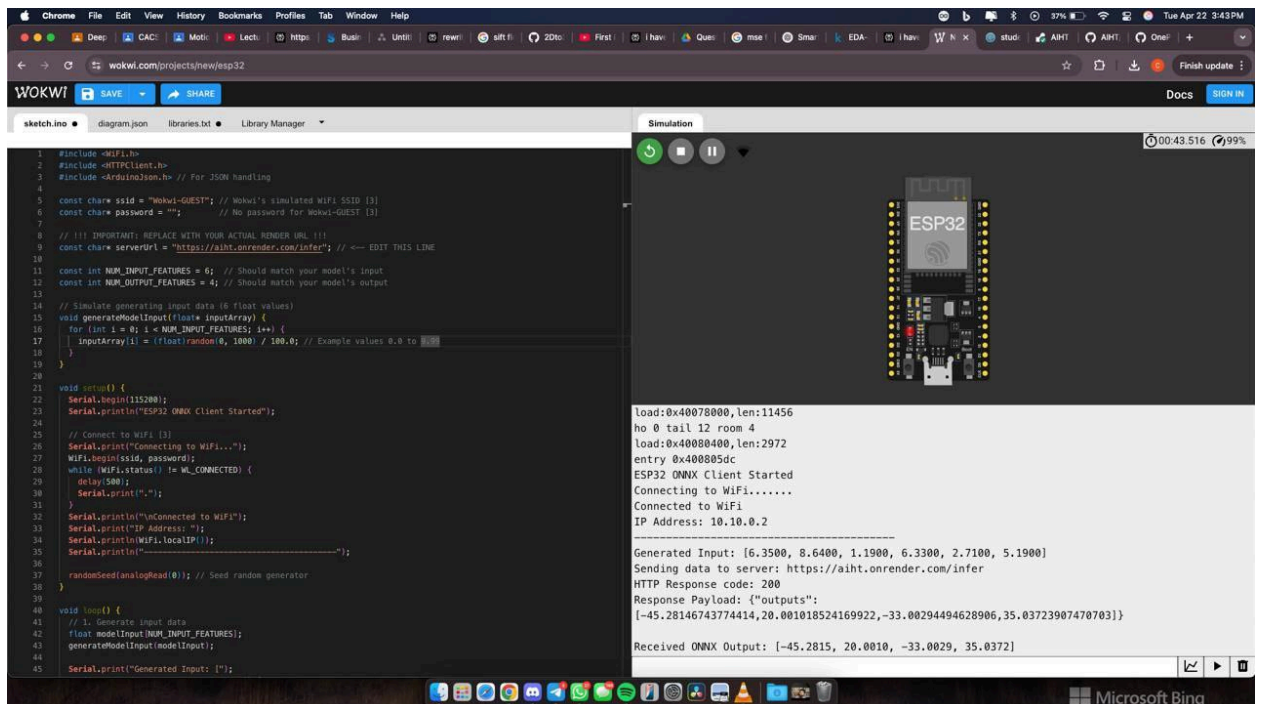
        input_array = np.array(input_values, dtype=np.float32).reshape(EXPECTED_INPUT_SHAPE)
        feeds = {INPUT_NAME: input_array}
        results = session.run([OUTPUT_NAME], feeds)
        output_data = results[0].flatten().tolist()

        log_to_csv(input_values, output_data)

        print(f"Received input: {input_values}, Produced output: {output_data}")
        return flask.jsonify({"outputs": output_data})

    except Exception as e:
        print(f"Error during inference: {e}")
        return flask.jsonify({"error": f"Inference error: {str(e)}"}), 500

```



The screenshot displays a web browser window with the Wokwi website. The main content area shows a project titled "sketch.ino" with a "diagram.json" file. The code editor on the left contains the following Arduino sketch:

```

1 #include <WiFi.h>
2 #include <HTTPClient.h>
3 #include <ArduinoJson.h> // For JSON handling
4
5 const char ssid = "Wokwi-GUEST"; // Wokwi's simulated WiFi SSID (3)
6 const char password = ""; // No password for Wokwi-GUEST (3)
7
8 // !!! IMPORTANT: REPLACE WITH YOUR ACTUAL RENDER URL !!!
9 const char serverUrl = "https://aiht.onrender.com/infer"; // <-- EDIT THIS LINE
10
11 const int NUM_INPUT_FEATURES = 6; // Should match your model's input
12 const int NUM_OUTPUT_FEATURES = 4; // Should match your model's output
13
14 // Simulate generating input data (6 float values)
15 void generateModelInput(float* inputArray) {
16     for (int i = 0; i < NUM_INPUT_FEATURES; i++) {
17         inputArray[i] = (float)random(0, 1000) / 100.0; // Example values 0.0 to 9.99
18     }
19 }
20
21 void setup() {
22     Serial.begin(115200);
23     Serial.println("ESP32 ONNX Client Started");
24
25     // Connect to WiFi (3)
26     Serial.println("Connecting to WiFi...");
27     WiFi.begin(ssid, password);
28     while (WiFi.status() != WL_CONNECTED) {
29         delay(500);
30         Serial.print(".");
31     }
32     Serial.println("\nConnected to WiFi");
33     Serial.println("IP Address: ");
34     Serial.println(WiFi.localIP());
35     Serial.println("-----");
36
37     randomSeed(analogRead(0)); // Seed random generator
38 }
39
40 void loop() {
41     // 1. Generate input data
42     float modelInput[NUM_INPUT_FEATURES];
43     generateModelInput(modelInput);
44
45     Serial.println("Generated Input: [");

```

The right pane shows the simulation interface with a 3D model of an ESP32. The terminal window displays the following output:

```

load:0x40070000,len:11456
ho 0 tail 12 room 4
load:0x40080400,len:2972
entry 0x400805dc
ESP32 ONNX Client Started
Connecting to WiFi.....
Connected to WiFi
IP Address: 10.10.0.2

Generated Input: [6.3500, 8.6400, 1.1900, 6.3300, 2.7100, 5.1900]
Sending data to server: https://aiht.onrender.com/infer
HTTP Response code: 200
Response Payload: {"outputs":
[-45.28146743774414,20.001018524169922,-33.00294494628906,35.03723907470703]}

Received ONNX Output: [-45.2815, 20.0010, -33.0029, 35.0372]

```